

# SQL Server 2022

Monday, April 24, 2023 8:13 PM

[SQL Server 2022 Blogging Series Archives - Microsoft SQL Server Blog](#)



## Industry-proven database engine

Scalability, management and availability

### Scalability



Buffer pool parallel scan

System page latch  
concurrency enhancements

Ordered clustered columnstore index

Improved columnstore segment  
elimination

In-memory OLTP enhancements

### Management



Query Store Improvements

Shrink database WAIT\_AT\_LOW\_PRIORITY

XML Compression

Backup and restore to S3-compatible  
object storage

Improved snapshot backup support

Improved Azure integration

### Availability



Always On Availability Group  
improvements

Link to Azure SQL Managed Instance

Multi-write replication (LWW)

Intel® QuickAssist Technology  
backup compression

Accelerated Database Recovery (ADR)

## SQL Server 2022

## Storage Engine Scalability and Management



System page latch  
concurrency  
enhancements



Buffer Pool Parallel  
Scan



In Memory OLTP  
Enhancements



Clustered Columnstore  
Indexing  
Improvements



Query Store  
enhancements



XML Compression



Shrink Database with  
Low Priority

# tempdb Performance Critical for Scalability

## SQL Server 2016 Improvements

- Setup experience has improved
- Trace Flag 1117 and 1118 are no longer required

## SQL Server 2017 Improvements

- Round robin PFS
- Improvements to optimistic latching

## SQL Server 2019 Improvements

- Memory-optimized tempdb metadata
- Concurrent PFS updates







## SQL Server 2022 Improvements

- System page latch concurrency enhancements
- Concurrent Global Allocation Map (GAM) Updates
- Concurrent Shared Global Allocation Map (SGAM) Updates

```
--GET OBJECT INFO FROM PAGE RESOURCE SQL 2019
-- RESOLVE WAIT RESOURCE TO OBJECT INFORMATION
-- NEW FUNCTIONS AVAILABLE IN SQL SERVER 2019
USE master GO
SELECT
er.session_id, er.wait_type, er wait_resource,
OBJECT_NAME(page_info.[object_id],page_info.database_id) as [object_name],
er.blockingsessionid,er.command,
SUBSTRING(st.text, (er.statement_start_offset/2) * 1,
((CASE er.statement_end_offset WHEN -1 THEN DATALENGTH(st.text)
ELSE er.statement_end_offset
END   er.statement_start_offset)/2) + 1) AS statement_text,
page_info.database_id,page_info.[file_id], page_info.page_id, page_info.[object_id],
page_info.index_id, page_info.page_type_desc FROM sys.dm_exec_requests AS er
CROSS APPLY sys.dm_exec_sql_text(er.sql_handle) AS st CROSS APPLY sys.fn_PageResCracker
(er.page_resource) AS r
CROSS APPLY sys.dm_db_page_info(r.[db_id], r.[file_id], r.page_id, 'DETAILED') AS page_info WHERE
er.wait_type like '%page%'
```

## Current Buffer Pool Scan Operations



- Operations that scan the buffer pool can be slow, especially on large memory machines, such as:
  - Creating a new databases 
  - File drop operations 
  - Backup/restore operations 
  - Always On failover events 
  - DBCC CHECKDB 
  - Log restore operations 
  - Internal operations (e.g., checkpoint)

# SQL Server 2022 - Buffer Pool Parallel Scan



- Buffer Pool Scans are parallelizing by utilizing multiple cores
- Benefits to both small and large databases on large-memory machines
- Improvement adds Buffer Pool scan diagnostics and telemetry for supportability and insights.
  - Long Buffer Pool scans will be visible by the ERRORLOG
  - Extended events will capture scan start/complete, errors, FlushCache, etc.
- **Customers running mission critical OLTP, and data warehouse environments will see the most benefit**

operations on large memory machines by  
utilizing multiple CPU cores

## References:

Improve scalability with Buffer Pool Parallel Scan in  
SQL Server 2022  
<https://aka.ms/bpps>

## Buffer Pool Parallel Scan Diagnostics

### Extended Events

event name	channel	timestamp	name	database_id	elapsed_time_ms	elapsed_time_sec	command	operation	scanned_buffers	total_buffers	sql_text
buffer_pool_scan_start	Debug	2020-04-21 08:30:03.7232160	buffer_pool_scan_complete	9	14249	14	CREATE DATABASE	FlushCache	115	204640229	create database testdb1
buffer_pool_scan_complete	Analytic	2020-04-21 08:35:57.7701836	buffer_pool_scan_complete	9	14277	14	DBCC	FlushCache	321	204640229	DBCC FLUSH(database, 'testdb1')
buffer_pool_scan_task_start	Debug	2020-04-21 08:36:59.6869784	buffer_pool_scan_complete	9	11422	11	DBCC	MarkBufferCopyOnWrite	321	204640229	dbcc checkdb(testdb1)
buffer_pool_scan_task_complete	Debug	2020-04-21 08:37:25.4003359	buffer_pool_scan_complete	10	14344	14	DBCC	FlushCache	26	204640229	dbcc checkdb(testdb1)
buffer_pool_scan_task_error	Debug	2020-04-21 08:37:48.5725842	buffer_pool_scan_complete	9	11545	11	DBCC CHECKCATALOG	MarkBufferCopyOnWrite	329	204640229	dbcc checkdb(testdb1)
buffer_pool_scan_stats	Analytic	2020-04-21 08:38:00.0099010	buffer_pool_scan_complete	10	11429	11	DBCC CHECKCATALOG	RemoveDatabase	383	204640229	dbcc checkdb(testdb1)
buffer_pool_flush_cache_start	Debug	2020-04-21 08:38:53.0208060	buffer_pool_scan_complete	9	11424	11	DBCC TABLE CHECK	MarkBufferCopyOnWrite	329	204640229	dbcc checktable(testdb1 db=1)
buffer_pool_flush_cache_status	Debug	2020-04-21 08:39:18.7317759	buffer_pool_scan_complete	10	14220	14	DBCC TABLE CHECK	FlushCache	26	204640229	dbcc checktable(testdb1 db=1)
buffer_pool_flush_cache_end	Debug	2020-04-21 08:39:41.6384020	buffer_pool_scan_complete	9	11417	11	DBCC	MarkBufferCopyOnWrite	329	204640229	dbcc checktable(testdb1 db=1)
buffer_pool_flush_cache_error	Analytic	2020-04-21 08:39:53.0759504	buffer_pool_scan_complete	10	11423	11	DBCC	RemoveDatabase	157	204640229	dbcc checktable(testdb1 db=1)
buffer_pool_flush_cache_stats	Debug	2020-04-21 08:42:33.3964545	buffer_pool_scan_complete	9	11461	11	ALTER DATABASE	RemoveDatabaseByFile	352	204640229	alter database testdb1 remove file data2
buffer_pool_flush_cache_end	Analytic	2020-04-21 08:43:46.0099952	buffer_pool_scan_complete	9	14232	14	BACKUP DATABASE	FlushCache	352	204640229	backup database testdb1 to disk='nul'
buffer_pool_flush_cache_status	Debug	2020-04-21 08:45:12.0203296	buffer_pool_scan_complete	10	14280	14	CREATE DATABASE	FlushCache	115	204640229	create database testdb3
buffer_pool_flush_cache_end	Analytic	2020-04-21 08:45:23.0216234	buffer_pool_scan_complete	10	14349	14	BACKUP DATABASE	FlushCache	306	204640229	create database testdb3
buffer_pool_flush_cache_status	Debug	2020-04-21 08:47:10.3404096	buffer_pool_scan_complete	10	11814	11	DROP DATABASE	RemoveDatabase	306	204640229	create database testdb3
buffer_pool_flush_cache_end	Debug	2020-04-21 08:47:36.5403815	buffer_pool_scan_complete	10	14470	14	RESTORE DATABASE	FlushCache	4	204640229	create database testdb3
buffer_pool_flush_cache_status	Debug	2020-04-21 08:48:13.7775913	buffer_pool_scan_complete	10	14223	14	RESTORE LOG	FlushCache	4	204640229	create database testdb3
buffer_pool_flush_cache_end	Debug	2020-04-21 08:48:51.0420029	buffer_pool_scan_complete	10	14344	14	RESTORE LOG	FlushCache	5	204640229	create database testdb3

## In-memory Feature Family

- In-Memory OLTP (SQL 2014+/2022)
- Memory-Optimized tempdb Metadata (SQL 2019/2022)
- Memory Protection Keys (SQL 2019/2022)
- Persisted Log Buffer ("Tail of Log Caching") (SQL 2016+/2019)
- Data file "Enlightenment" (SQL 2019/2022)
- Hybrid Buffer Pool (SQL 2019/2022)
  - Read caching (SQL 2019)
  - Direct write (SQL 2022)

## In-memory OLTP memory enhancements



- Improved supportability with insights on memory usage related to memory consumers  
`sys.dm_xtp_system_memory_consumers`
- Improved the ability to manage and make memory shrinkable
- Release as much memory as possible in case of memory pressure or on demand
- Provide a method to release unused memory on demand  
`sys.sp_xtp_force_gc`

## Columnstore Index Improvements



- Ordered Clustered ColumnStore Index - Ordered (CCI) sorts the existing data in memory before the index builder compresses the data into index segments
  - More efficient segment elimination
  - Reduced number of segments to read from disk
- Improved ColumnStore segment elimination - Beginning in SQL Server 2022, segment elimination capabilities extend to string, binary, GUID data types, and the datetimeoffset data types
- Improved optimization - SQL Server 2022 leverages new hardware capabilities including the Advanced Vector Extension (AVX) 512 extension to improve batch mode operations.

## XML Compression



- XML data type is commonly used to store unstructured data
- Data compression only applies to in-row scenarios (row, page)
- XML Compression will compress the XML data type in Azure SQL and SQL Server 2022
- XML Compression can be specified during CREATE and ALTER of TABLE and INDEX statements
- [sp\\_estimate\\_data\\_compression\\_savings](#) will be expanded to estimate XML savings

```
ALTER TABLE Sales.StoreBIGXMLCopy REBUILD PARTITION = ALL  
WITH (DATA_COMPRESSION = PAGE, XML_COMPRESSION = ON)
```

## Shrink Database with Low Priority

- Customers often need to reclaim data space
- Common for hosted environments (new database per customer)
- Shrink Database operations can cause concurrency issues
- Shrink Database WLP addresses this problem by waiting with less restrictive locking
- Similar to ALTER INDEX WAIT\_AT\_LOW\_PRIORITY

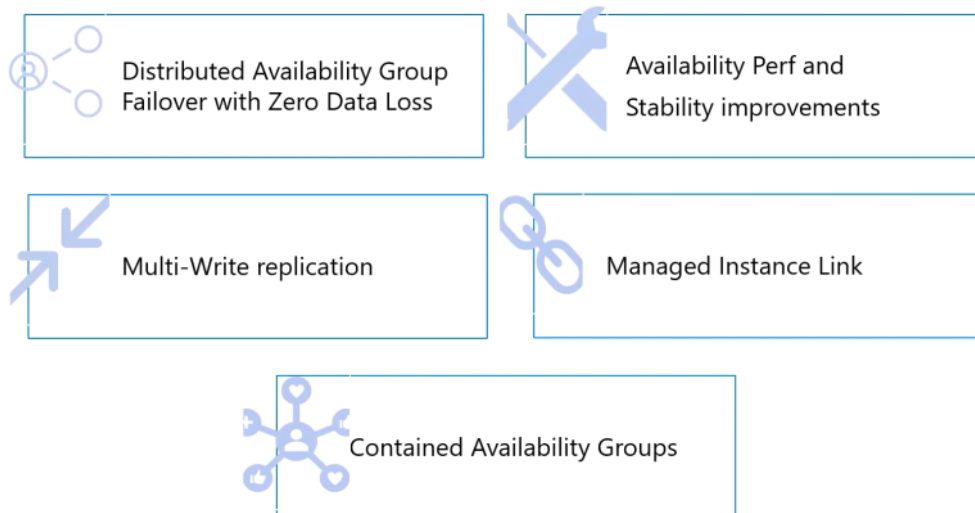
```
DBCC SHRINKDATABASE (2, 20, NOTTRUNCATE)  
WITH WAIT_AT_LOW_PRIORITY (ABORT_AFTER_WAIT = SELF))
```

```
DBCC SHRINKFILE (5, 7) WITH WAIT_AT_LOW_PRIORITY  
(ABORT_AFTER_WAIT = BLOCKERS))
```

### SQL Server 2022

Press Esc to exit full screen

### High Availability New Feature Overview



## Distributed Availability Groups

- Intended for remote locations
- Resolves quorum issues between sites
- Improves WAN network efficiency
- Enables mixed Windows versions

# Distributed Availability Groups

- Challenges:
  - No central cluster manager
  - No coordination to synchronize failovers
  - **Very difficult to guarantee zero data loss in failovers.**

## Zero Data loss Distributed AG Failover

- Add the ability to apply `REQUIRED_SYNCHRONIZED_SECONDARIES_TO_COMMIT` at the Distributed AG level
- Ensures that no transaction is partially committed when failover happens.
- Takes a very complex error-prone operation and makes it simple and success-prone.

## What is link and why use it?

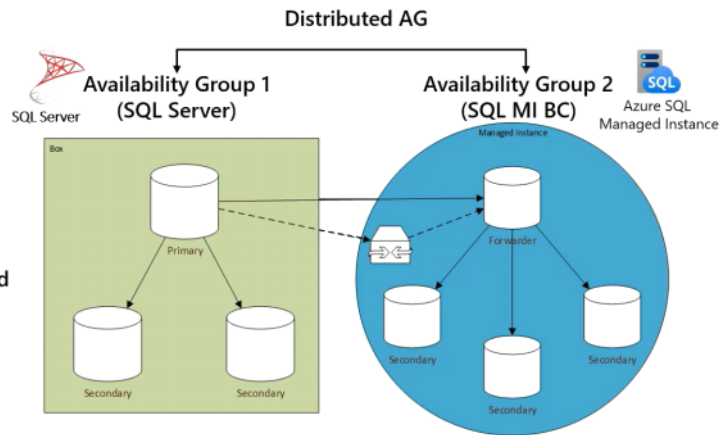
- A common capability in the SQL Server and SQL Managed Instance that enables rich **hybrid connectivity scenarios** between the two.
- Empowers you to immerse into Azure at your **own pace**.
- **Modernize** in Azure while keep running on your **existing SQL Server**.
- When and if ready, **migrate** to Azure at your own pace with **minimum downtime**.
- Use Managed Instance as a **Disaster Recovery (DR)** site (SQL Server 2022 only)\*



# Link is based on distributed Availability Groups

## Distributed Availability Groups

- Leveraging the proven Availability Group technology stack
- Connection between SQL Server and SQL MI is bridged using distributed AG
- Distributed AG spans two separate Availability Groups, across domains and across platforms
- WSFC is not required SQL Server 2017, 2019 and 2022.
- Single-node WSFC is required for SQL Server 2016.
- You can use mixed versions of SQL Servers to connect to SQL MI.



## AG Perf and Stability improvements

- Two TCP channels for log replication in Distributed AG environments
  - Helps compensate for high latency
  - Improves throughput
- Redo Thread Pool enhancements
  - Current: Statically assign threads to databases as they start up (100 threads/DB) Remaining DBs get one thread each.
  - NEW: Global thread pool for redo work. Any database can post a work item to the pool and get a thread in FIFO order. Makes for more fair and efficient usage of execution threads.
- AlwaysOn\_Health : Capture sp\_server\_diagnostics XEvent when STATE=3 (ERROR) To Diagnose HADR Health Events
- DCR: Add REVERTING progress to errorlog Just like SQL Already Reports Recover Progress
- DCR: Improve AlwaysOn\_health and system\_health logs, add XEvents that are helpful to understand database state, recovery and reverting progress
- Introduce unfair mutex into UCS session registration to improve AG data movement performance under multi-db scenario

## Multi-write replication

Multi-master writes for users across multiple locations

- Globally distributed database replicas for geo-localized writes
- Enhanced conflict detection for inserts and updates with Last Writer Wins (LWW) capabilities
- Ensures the last update is persisted across all replicas based on the UTC time of the operation.



## Contained Availability Groups

### Problem Statement

- Objects or settings related to an AG do not get replicated across the instances in the AG
  - Users/Logins
  - Permissions
  - Jobs
  - Settings
- Applications need a consistent execution context to function properly
- Workarounds involve complex and fragile scripts to force synchronization or manual processes

### Goals

- Provide a unified experience across all instances in the Contained AG
- Nearly all operations within the AG require no syntax change
  - The exception being operations specifically impacting a contained AG
- Preserve the operations and management of the individual instances in the Contained AG
- Enable multiple Contained AGs to coexist on the same set of instances without interference



# Contained AG Syntax

(Only the syntax in **RED** is new)

```
CREATE AVAILABILITY GROUP MyAg
    WITH (CONTAINED, [REUSE_SYSTEM_DATABASES,] AUTOMATED_BACKUP_PREFERENCE =
SECONDARY, FAILURE_CONDITION_LEVEL = 3, HEALTH_CHECK_TIMEOUT = 600000 )
    FOR DATABASE ThisDatabase, ThatDatabase
    REPLICAS ON
        'COMPUTER01' WITH ( ENDPOINT_URL = 'TCP://COMPUTER01:5022',
            AVAILABILITY_MODE = SYNCHRONOUS_COMMIT),
        'COMPUTER02' WITH ( ENDPOINT_URL = 'TCP://COMPUTER02:5022',
            AVAILABILITY_MODE = SYNCHRONOUS_COMMIT))
```

```
ALTER AVAILABILITY GROUP [MyAg] ADD LISTENER 'MyAgListenerIPv6' ( WITH IP (
('2001:db88:f0:f00f::cf3c'),('2001:4898:e0:f213::4ce2') ) , PORT = 60173 ); GO
```

## Managing the Contained AG

- **NOTE: Contained Availability Groups do NOT represent a security boundary.**
- To manage or interact with the Contained AG, connect to the contained AG:
  - Using the Contained AG Listener
  - Specifying a database within the Contained AG
    - USE MyAG\_MSDB
- Database names must be unique across all instances

## System Databases

- We create <AG-Name>\_Master and <AG-Name>\_MSDB in each Contained AG.
- They are initially empty but contain the subset of the instance master which we replicate. (Logins, users, permissions, jobs, etc.)
- They are automatically seeded to new AG replicas

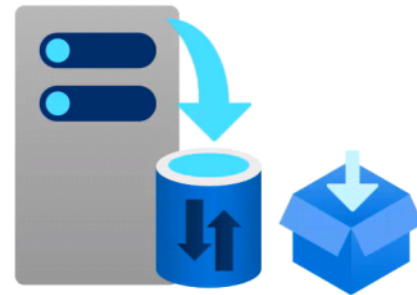
[SQL Server Contained Availability Groups Configuration \(mssqltips.com\)](https://www.mssqltips.com/sqlserver/2012/02/sql-server-contained-availability-groups-configuration/)

SQL Server 2022 - Backup  
Compression and Offloading with  
Intel® QuickAssist Technology (QAT)

# Hardware acceleration for backup compression

## Why compress backups?

- Reduce capacity
- Improve performance
  - Faster backup
  - Faster restore
- Existing backup compression products
  - Software based compression adds CPU overhead
  - Backup speed is variable
  - Depends on available CPU resources
  - More likely to impact workload



## Intel® QuickAssist Technology (QAT) compression

- Intel QuickAssist Technology improves performance across applications, including symmetric encryption and authentication, asymmetric encryption, digital signatures, RSA, DH, and ECC, and lossless data compression.
- [Intel QuickAssist Technology \(QAT\)](#) provides a compression technology that integrates hardware acceleration of compute-intensive workloads on Intel platforms.

### Database Backup Improvements for SQL Server 2022 with Intel QuickAssist Technology:



**Up to 2.3x**  
speedup in database  
backup time<sup>1</sup>



**Up to 6%**  
reduction in backup  
storage capacity<sup>1</sup>

# SQL Server 2022 and Intel® QAT

## Enabling Hardware Offloading

- Leverages Intel® QuickAssist Technology (Intel® QAT) for improved backup performance
- Free-up processor cycles by offloading backup compression
- Reduce demands on processor
- Improves backup speed

Enable the hardware offload feature for SQL Server 2022:

```
sp_configure 'hardware offload enabled', 1
GO
RECONFIGURE
GO
```



Enable Intel® QuickAssist Technology (QAT) hardware mode:

```
ALTER SERVER CONFIGURATION
SET HARDWARE_OFFLOAD = ON
(ACCELERATOR = QAT)
```

Enable Intel® QuickAssist Technology (QAT) software mode:

```
ALTER SERVER CONFIGURATION
SET HARDWARE_OFFLOAD = ON
(ACCELERATOR = QAT, MODE = SOFTWARE)
```

Backup database with QAT compression:

```
BACKUP DATABASE [TicketReservations]
TO DISK = 'D:\backups\QAT-DEFLATE.bak'
WITH FORMAT, COMPRESSION (ALGORITHM = QAT_DEFLATE)
GO
```

Backup database with the default compression:

```
BACKUP DATABASE [TicketReservations]
TO DISK = 'D:\backups\MS-XPRESS.bak'
WITH FORMAT
GO
```

## SQL Server 2022 and Intel® QAT – Restore Headeronly

- Restore headeronly will display the compression algorithm
- Running Restore headeronly requires that the Intel® QuickAssist Technology (Intel® QAT) drivers are available

RESTORE HEADERONLY

FROM DISK = 'D:\backups\QAT-Deflate.bak'

GO

KeyAlgorithm	EncryptorThumbprint	EncryptorType	LastValidRestor...	TimeZone	CompressionAlgorithm
NULL	NULL	NULL	NULL	-1	MS-XPRESS
KeyAlgorithm	EncryptorThumbprint	EncryptorType	LastValidResto...	TimeZone	CompressionAlgorithm
NULL	NULL	NULL	NULL	0	QAT-DEFLATE

## SQL Server 2022 and Intel® QAT – Restore Database

- SQL Server backups compressed using QAT\_DEFLATE will support all T-SQL RESTORE operations
- Intel® QAT compressed backups require the server running the RESTORE DATABASE command to have the Intel® QAT drivers installed

Restore database command:

RESTORE DATABASE [testdb1]

FROM DISK = 'D:\backups\QAT-Deflate.bak'

GO



# SQL Server 2022 and Intel® QAT - Troubleshooting

## sys.dm\_server\_accelerator\_status

- The SQL Server sys.dm\_server\_accelerator\_status dynamic management view can be used to verify the configuration state
- The mode description reason is a valuable for troubleshooting
- Configured libraries and driver versions are visible

The DMV below will expose offloading potentials and the configuration, one row per accelerator:



```
SELECT * FROM sys.dm_server_accelerator_status
GO
```

## SQL Server 2022 and Intel® QAT - Troubleshooting

- Intel® QuickAssist Technology Landing Page  
<https://developer.intel.com/quickassist>
- SQL Server error log will show if hardware has been detected along with the accelerator mode

	LogDate	ProcessInfo	Text
1	2022-08-16 23:35:05.420	Server	Detected Intel(R) QuickAssist Compression Library QATZip: 1.8.0.10.
2	2022-08-16 23:35:05.420	Server	Detected Intel(R) QuickAssist Kernel Driver icp_qat: 3.2.0.9.
3	2022-08-16 23:35:05.420	Server	Detected Intel(R) Storage Acceleration Library ISA-L: 2.30.0.0.
4	2022-08-16 23:35:05.420	Server	Intel(R) QuickAssist Technology (QAT) initialization succeeded.
5	2022-08-16 23:35:05.420	Server	Intel(R) QuickAssist Technology (QAT) hardware detected on the system.
6	2022-08-16 23:35:05.420	Server	Intel(R) QuickAssist Technology (QAT) will be used in hardware mode.
7	2022-08-16 23:35:05.420	Server	Intel(R) QuickAssist Technology (QAT) session successfully created.

- [Some Intel QAT Backup Compression Results](#)
- [How to Enable Intel QAT Backup Compression in SQL Server 2022](#)
- [How to Install the Intel QAT Driver](#)

--#8. Using msdb and the compression\_algorithm via backupset for backup performance history

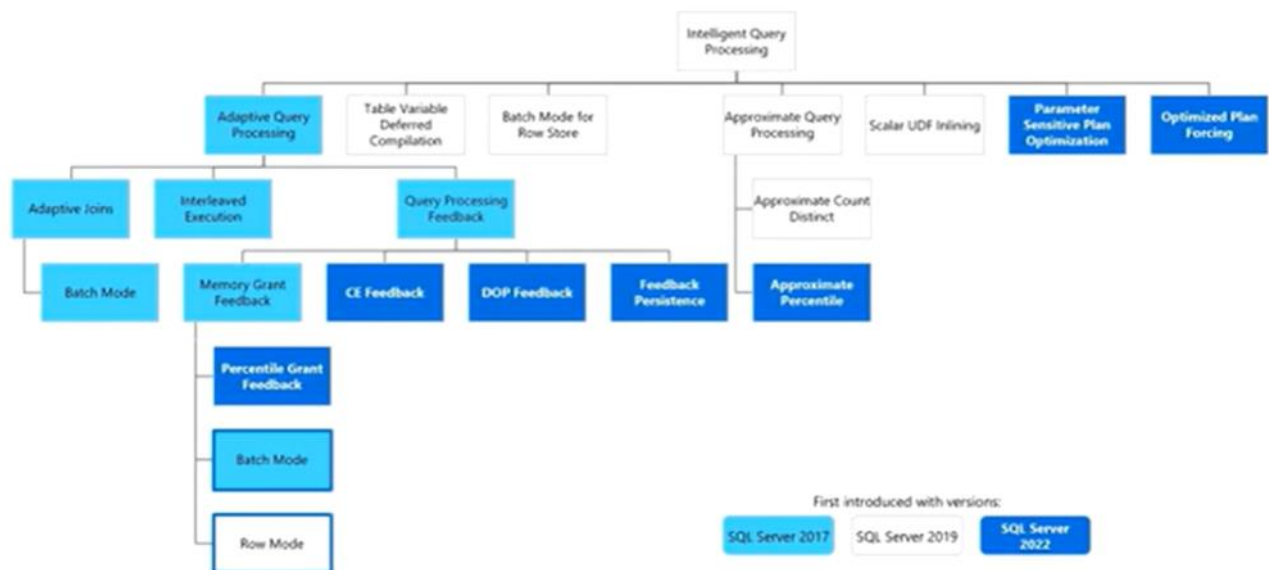
```
SELECT bs.databasesname,  
backuptype = CASE  
WHEN bs.type = 'D'  
AND bs.is_copy_only = 0 THEN 'Full Database'  
WHEN bs.type = 'D'  
AND bs.is_copy_only = 1 THEN 'Full Copy-Only Database'  
WHEN bs.type = 'I' THEN 'Differential database backup'  
WHEN bs.type = 'L' THEN 'Transaction Log'  
WHEN bs.type = 'F' THEN 'File or filegroup'  
WHEN bs.type = 'G' THEN 'Differential file'  
WHEN bs.type = 'P' THEN 'Partial'
```

```

WHEN bs.type = 'Q' THEN 'Differential partial'
END + ' Backup',
CASE bf.device_type
WHEN 2 THEN 'Disk'
WHEN 5 THEN 'Tape'
WHEN 7 THEN 'Virtual device'
WHEN 9 THEN 'Azure Storage'
WHEN 105 THEN 'A permanent backup device'
ELSE 'Other Device'
END AS DeviceType,
bms.software_name AS backupsoftware,
bs.recovery_model,
bs.compatibility_level,
BackupStartDate = bs.Backup_Start_Date,
BackupFinishDate = bs.Backup_Finish_Date, DATEDIFF(ms, backup_start_date, backup_finish_date) AS
ms_taken,
LatestBackupLocation = bf.physical_device_name,
backup_size_mb = CONVERT(decimal(10, 2), bs.backup_size/1024./1024.),
compressed_backup_size_mb = CONVERT(decimal(10, 2), bs.compressed_backup_size/1024./1024.),
(CONVERT(decimal(10, 2), bs.backup_size/1024./1024.) - (CONVERT(decimal(10, 2),
bs.compressed_backup_size/1024./1024.)) as compression_savings, bs.compressionalgorithm
FROM msdb.dbo.backupset bs
LEFT OUTER JOIN msdb.dbo.backupmediafamily bf ON bs.[media_set_id] = bf.[media_set_id]
INNER JOIN msdb.dbo.backupmediaset bms ON bs.[media_set_id] = bms.[media_set_id]
WHERE bs.backup_start_date > DATEADD(DAY, -72, sysdatetime()) --only look at last 12 hours
ORDER BY bs.Backup_Start_Date DESC, bs.databasesname ASC

```

## IQP Family Tree



Query Store - Custom capture policies

```

ALTER OATABASE [QueryStoreDB]
SET QUERY_STORE = ON
(
OPERATION_MODE = READ_WRITE,
CLEANUP_POLICY = ( STALE_QUERY_THRESHOLD_DAYS = 90 ),
DATA_FLUSH_INTERVAL_SECONDS = 900,
MAX_STORAGE_SIZE_MB = 1000, INTERVAL_LENGTH_MINUTES = 60, SIZE_BASED_CLEANUP_MODE =
AUTO, MAX_PLANS_PER_QUERY = 200, WAIT_STATS_CAPTURE_MODE = ON,
QUERY_CAPTURE_MODE = CUSTOM,
QUERY_CAPTURE_POLICY = (
STALE_CAPTURE_POLICY_THRESHOLD = 24 HOURS, EXECUTION_COUNT = 30,
TOTAL_COMPILE_CPU_TIME_MS = 1000, TOTAL_EXECUTION_CPU_TIME_MS = 100 )
)

```



);

## What are Query hints?



- Ideally the Query Optimizer selects an optimal execution plan
- Developers and DBAs often need to influence plan behavior but historically had few options outside of modifying application code
- Query hints are specified via the OPTION clause to influence the behavior of operators in a statement

### Example:

```
SELECT COUNT(DISTINCT [WWI Order ID])  
FROM [Fact].[OrderHistoryExtended]  
OPTION (USE HINT('DISALLOW_BATCH_MODE'), RECOMPILE);
```

## Applying Query hints Today

- Query hints help provide solutions to various performance related issues – but they do require a rewrite of the original query
- DBAs often cannot make changes directly to T-SQL code
  - T-SQL hard-coded into application
  - T-SQL automatically generated by the application
- DBA may have to rely on plan guides
  - Common feedback: “plan guides are complex to use and manage”

# Introducing Query Store hints

Query Store hints provides a method to shape query plans without changing code

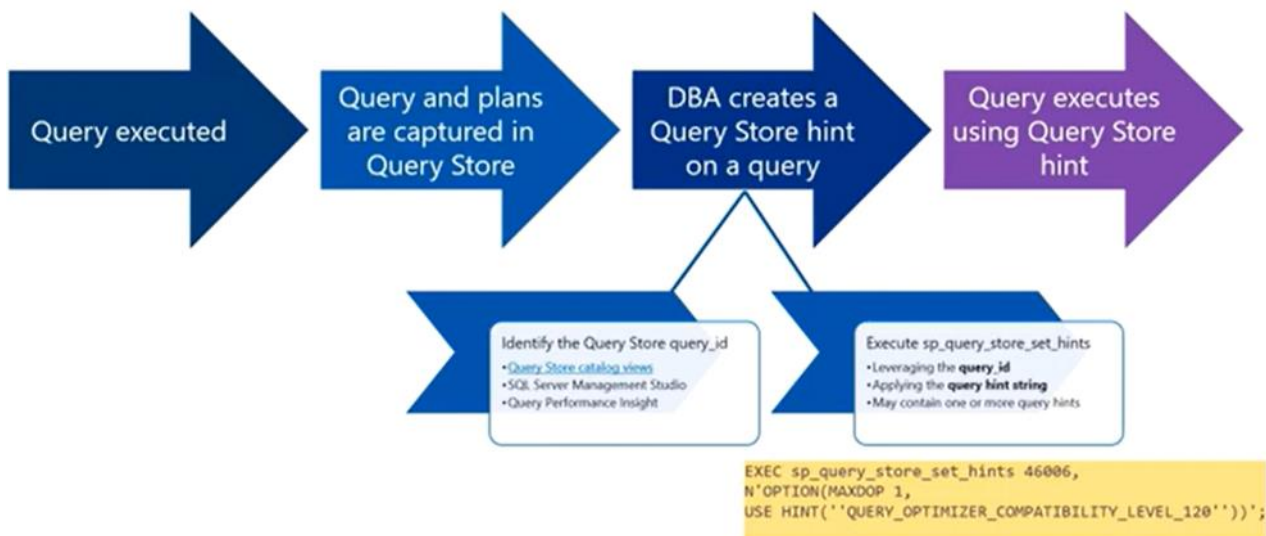
Query Store hints extends the power of Query Store

Query Store hints are persisted, surviving restarts

Override other hard-coded statement level hints and plan guides

Queries will always execute as opposing Query Store hints will be ignored

## Using Query Store hints



[Query Store hints in SQL Server 2022 - Microsoft SQL Server Blog](#)

# Feedback Workflow



User executes a query



If query qualifies for feedback, possible changes are suggested to the system



Changes are tried, and if they improve query performance, are verified and persisted

## Cardinality Estimation – What is the model?

- Basic set of assumptions that determine how we interpret statistics and predicates used in the query
- There has historically been one “model” for CE, but not all queries perform best with this model
- CE Feedback allows us to modify the underlying assumptions for a specific query – changing the model for that query and improving performance

## What is CE Feedback?

- Start with default model assumption
- Modify model assumption based on observations about query execution
- Validate that new model assumption is better

## Model assumption 1: Uniformity

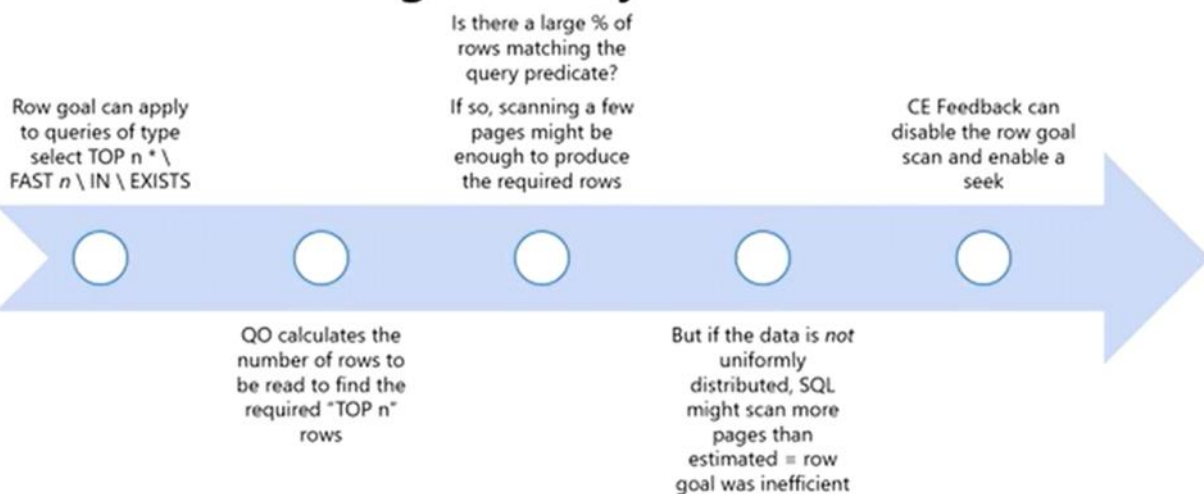


Uniformity is used when interpreting data from on-disk histograms



Data within histogram buckets is assumed to be uniform

## CE Feedback – Row goal Analysis



## Model assumption 2: Independence (or not...)

- Calculating the selectivity of a conjunction of predicates that are independent is done by multiplying their individual selectivities
- Quite often, predicates are not independent
- ...and sometimes they are not completely correlated, either